

▼ Reading the Data and Observations on Dataset

```
#Importing necessary libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

cd sample_data

[Errno 2] No such file or directory: 'sample_data'
/content/sample_data

import pandas as pd
data = pd.read_excel('customer_churn_large_dataset.xlsx')

data.head(5)
```

	CustomerID	Name	Age	Gender	Location	Subscription_Length_Mont
0	1	Customer_1	63	Male	Los Angeles	
1	2	Customer_2	62	Female	New York	
2	3	Customer_3	24	Female	Los Angeles	

```
print(len(data))

100000

# Display basic information about the dataset
print(data.info())

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   CustomerID            100000 non-null  int64
1   Name                  100000 non-null  object
2   Age                   100000 non-null  int64
3   Gender                100000 non-null  object
4   Location               100000 non-null  object
5   Subscription_Length_Months  100000 non-null  int64
6   Monthly_Bill           100000 non-null  float64
7   Total_Usage_GB         100000 non-null  int64
8   Churn                  100000 non-null  int64
dtypes: float64(1), int64(5), object(3)
memory usage: 6.9+ MB
None

print(data.describe())
```

	CustomerID	Age	Subscription_Length_Months	
count	100000.000000	100000.000000	100000.000000	
mean	50000.500000	44.027020	12.490100	
std	28867.657797	15.280283	6.926461	
min	1.000000	18.000000	1.000000	
25%	25000.750000	31.000000	6.000000	
50%	50000.500000	44.000000	12.000000	
75%	75000.250000	57.000000	19.000000	
max	100000.000000	70.000000	24.000000	
	Monthly_Bill	Total_Usage_GB	Churn	
count	100000.000000	100000.000000	100000.000000	
mean	65.053197	274.393650	0.497790	
std	20.230696	130.463063	0.499998	

```

min      30.000000      50.000000      0.000000
25%      47.540000     161.000000      0.000000
50%      65.010000     274.000000      0.000000
75%      82.640000     387.000000      1.000000
max     100.000000     500.000000      1.000000

```

```
data.isna().sum()
```

```

CustomerID      0
Name            0
Age            0
Gender          0
Location        0
Subscription_Length_Months  0
Monthly_Bill    0
Total_Usage_GB  0
Churn           0
dtype: int64

```

Great! There are no missing values.

▼ Data preprocessing and cleaning

Decided to remove the CustomerID, Name columns as they won't making any impact on model.

```
data.drop(["CustomerID"], inplace = True, axis = 1)
```

```
data.drop(["Name"], inplace = True, axis = 1)
```

```
data
```

	Age	Gender	Location	Subscription_Length_Months	Monthly_Bill	Tc
0	63	Male	Los Angeles	17	73.36	
1	62	Female	New York	1	48.76	
2	24	Female	Los Angeles	5	85.47	
3	36	Female	Miami	3	97.94	
4	46	Female	Miami	19	58.14	
...	
99995	33	Male	Houston	23	55.13	
99996	62	Female	New York	19	61.65	
99997	64	Male	Chicago	17	96.11	

```

# Handling missing data
data.dropna(inplace=True)

```

```

# Handling outliers (you can use appropriate methods depending on the distribution of your data)
# For example, using z-score for numerical columns
from scipy.stats import zscore
z_scores = zscore(data[['Age', 'Subscription_Length_Months', 'Monthly_Bill', 'Total_Usage_GB']])
data = data[(z_scores < 3).all(axis=1)]

```

```
data
```

	Age	Gender	Location	Subscription_Length_Months	Monthly_Bill	Tc
0	63	Male	Los Angeles	17	73.36	
1	62	Female	New York	1	48.76	
2	24	Female	Los Angeles	5	85.47	
3	36	Female	Miami	3	97.94	
4	46	Female	Miami	19	58.14	
...	
99995	33	Male	Houston	23	55.13	
99996	62	Female	New York	19	61.65	
99997	64	Male	Chicago	17	96.11	

```
# Example: Remove outliers using IQR
Q1 = data['Total_Usage_GB'].quantile(0.25)
Q3 = data['Total_Usage_GB'].quantile(0.75)
IQR = Q3 - Q1
data = data[(data['Total_Usage_GB'] >= Q1 - 1.5 * IQR) & (data['Total_Usage_GB'] <= Q3 + 1.5 * IQR)]
```

```
# Let's say 'Gender' is a categorical column you want to encode
data= pd.get_dummies(data, columns=['Gender'], drop_first=True)
```

```
# Example: Min-Max Scaling
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
data[['Age', 'Monthly_Bill']] = scaler.fit_transform(data[['Age', 'Monthly_Bill']])
```

data

	Age	Location	Subscription_Length_Months	Monthly_Bill	Total
0	0.865385	Los Angeles	17	0.619429	
1	0.846154	New York	1	0.268000	
2	0.115385	Los Angeles	5	0.792429	
3	0.346154	Miami	3	0.970571	
4	0.538462	Miami	19	0.402000	
...	
99995	0.288462	Houston	23	0.359000	
99996	0.846154	New York	19	0.452143	
99997	0.884615	Chicago	17	0.944429	

▼ feature engineering

```
# Example: Create a feature for the ratio of Monthly_Bill to Total_Usage_GB
data['Bill_to_Usage_Ratio'] = data['Monthly_Bill'] / data['Total_Usage_GB']
```

```
# Example: Create an interaction feature
data['Subscription_Bill_Interaction'] = data['Subscription_Length_Months'] * data['Monthly_Bill']
```

```
# Example: Calculate the average Monthly_Bill by Location
average_bill_by_location = data.groupby('Location')['Monthly_Bill'].mean().reset_index()
average_bill_by_location.rename(columns={'Monthly_Bill': 'Avg_Monthly_Bill_Location'}, inplace=True)
data = data.merge(average_bill_by_location, on='Location', how='left')
```

```
# Example: If you know that certain Age ranges are more likely to churn, create an Age_Group feature
data['Age_Group'] = pd.cut(data['Age'], bins=[0, 25, 35, 50, 100], labels=[0, 1,2,3])
```

data

	Age	Location	Subscription_Length_Months	Monthly_Bill	Total
0	0.865385	Los Angeles	17	0.619429	
1	0.846154	New York	1	0.268000	
2	0.115385	Los Angeles	5	0.792429	
3	0.346154	Miami	3	0.970571	
4	0.538462	Miami	19	0.402000	
...
99995	0.288462	Houston	23	0.359000	
99996	0.846154	New York	19	0.452143	
99997	0.884615	Chicago	17	0.944429	
99998	0.634615	New York	20	0.275000	
99999	0.173077	Los Angeles	19	0.665286	

```
data = pd.read_excel('customer_churn_large_dataset.xlsx')
correlation_matrix = data.corr()
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', center=0)
plt.title('Correlation Heatmap')
plt.show()
```

```
<ipython-input-127-929b20b3ec89>:2: FutureWarning: The default value of nu
correlation_matrix = data.corr()
```



Customer churn is highly correlated with Age and Subscription_Lenght_Month

```
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder

# Encode categorical variables
encoder = LabelEncoder()
data['Gender'] = encoder.fit_transform(data['Gender_Male'])
data['Location'] = encoder.fit_transform(data['Location'])

# Split the data into features (X) and target (y)
X = data.drop('Churn', axis=1)
y = data['Churn']

# Split into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Feature engineering
X_train['Bill_to_Usage_Ratio'] = X_train['Monthly_Bill'] / X_train['Total_Usage_GB']
X_test['Bill_to_Usage_Ratio'] = X_test['Monthly_Bill'] / X_test['Total_Usage_GB']
```

X_train

	Age	Location	Subscription_Length_Months	Monthly_Bill	Total
75220	0.692308	4	5	0.778571	
48955	0.192308	4	24	0.743714	
44966	0.750000	0	12	0.318429	
13568	0.019231	1	19	0.036714	
92727	0.730769	3	8	0.050286	
...
6265	0.326923	3	21	0.533286	
54886	0.730769	0	13	0.791429	
76820	0.980769	1	2	0.660571	
860	0.711538	0	12	0.845571	
15795	0.153846	2	17	0.577286	

```
import pandas as pd
from sklearn.preprocessing import StandardScaler
```

```
# Apply feature scaling
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Print the first few rows of the scaled dataset
print(X_train_scaled)

[[ 6.53446686e-01  1.41703505e+00 -1.08272837e+00 ...  8.41406774e-01
  0.00000000e+00 -9.95534968e-01]
 [-1.04827574e+00  1.41703505e+00  1.66388226e+00 ...  8.41406774e-01
  0.00000000e+00  1.00448506e+00]
 [ 8.49799274e-01 -1.41806277e+00 -7.08191867e-02 ... -1.61937346e+00
  0.00000000e+00  1.00448506e+00]
 ...
 [ 1.63520963e+00 -7.09288318e-01 -1.51640373e+00 ...  5.15808394e-01
  0.00000000e+00  1.00448506e+00]
 [ 7.18897549e-01 -1.41806277e+00 -7.08191867e-02 ... -1.61937346e+00
  0.00000000e+00  1.00448506e+00]
 [-1.17917747e+00 -5.13861481e-04  6.51973084e-01 ...  9.61784624e-01
  0.00000000e+00 -9.95534968e-01]]
```

▼ Model selection and optimization.

```
import pandas as pd
from sklearn.impute import SimpleImputer
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score

from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

imputer = SimpleImputer(strategy='mean') # You can use other strategies as well

# Impute missing values in the feature matrix
X_imputed = imputer.fit_transform(X_train)
X_test_imputed= imputer.transform(X_test)
# Initialize the model
model = RandomForestClassifier(random_state=42)
# Train the model
model.fit(X_imputed, y_train)
```

```
▼ RandomForestClassifier
RandomForestClassifier(random_state=42)
```

```
y_pred=model.predict(X_test_imputed)
```

[+ Code](#)
[+ Text](#)

```
# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print(f"Accuracy: {accuracy:.2f}")
print(f"Precision: {precision:.2f}")
print(f"Recall: {recall:.2f}")
print(f"F1 Score: {f1:.2f}")
```

```
Accuracy: 0.50
Accuracy: 0.50
Precision: 0.50
Recall: 0.47
F1 Score: 0.48
```

```
from sklearn.model_selection import GridSearchCV

# Define hyperparameters to tune
param_grid = {
    'n_estimators': [10, 20, 30],
    'max_depth': [None, 10, 20],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}

# Initialize GridSearchCV
grid_search = GridSearchCV(estimator=model, param_grid=param_grid, cv=3, scoring='f1')

# Perform grid search
grid_search.fit(X_imputed, y_train)

# Get the best model
best_model = grid_search.best_estimator_

# Make predictions using the best model
y_pred_best = best_model.predict(X_test_imputed)

# Evaluate the best model
f1_best = f1_score(y_test, y_pred_best)
print(f"Best F1 Score: {f1_best:.2f}")
```

Best F1 Score: 0.49

